

[1] Dynamics of Implied Volatility Surfaces [Rama Cont and Jose da Fonseca, 2002]

- Link:  
[https://www.researchgate.net/publication/227624113\\_Dynamics\\_of\\_Implied\\_Volatility\\_Surfaces](https://www.researchgate.net/publication/227624113_Dynamics_of_Implied_Volatility_Surfaces)
- Short summary
  - The paper looks at how the implied volatility surface (IVS) of option prices on the SP500 and FTSE indices moves over time, not just its static smile structure shape.
  - The authors argue that option markets have their own sources of randomness beyond the underlying index, and that IV surfaces behave like a random surface driven by a few factors, which creates “vega risk” for option portfolios.
  - They propose an empirical factor model of the IVS based on a Karhunen–Loeve (PCA) decomposition of daily IV changes.
- Key theory features
  - Implied volatility surface as state variable
    - They argue it's better to model implied vol directly than local volatility, because IV is observable, directly linked to traded options, and familiar to practitioners.
  - Smile rules vs stochastic surface
    - They discuss “sticky moneyness” (surface constant in moneyness coordinates)

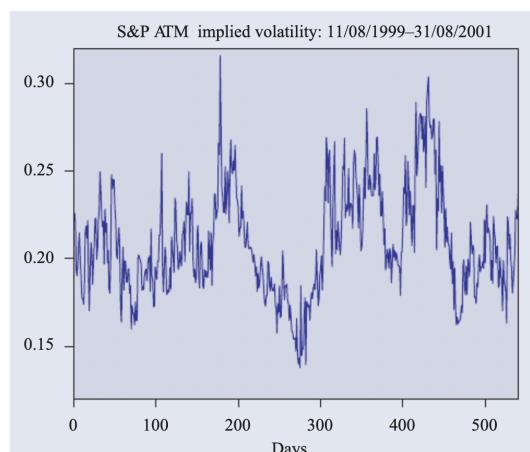
$$\forall(m, \tau), \quad I_{t+\Delta t}(m, \tau) = I_t(m, \tau)$$

and “sticky strike” (IV fixed for each (K,T))

$$\forall(K, T), \quad \sigma_{t+\Delta t}^{\text{BS}}(K, T) = \sigma_t^{\text{BS}}(K, T)$$

, which traders use as deterministic update rules for IVS.

- By showing large day-to-day variation in IV (e.g. ATM IV moving 15-40% on S&P in a few months



**Figure 2.** Evolution of at the money implied volatility for SP500 options, Aug. 1999–Aug. 2001.

), they argue these rules are too simplistic and can't capture real volatility risk.

- Implied volatility as a random surface
  - They model the (log) implied volatility surface as a stationary random field on the 2D domain (moneyness, time-to-maturity).
- Karhunen–Loeve decomposition - functional PCA of surfaces
  - They apply a Karhunen–Loève (KL) decomposition, i.e. PCA of the random surface:

$$U(\omega, \cdot) = \sum U_n(\omega) f_n(\cdot) = \sum U_n f_n.$$

where  $f_n(x)$  are deterministic eigen-surfaces (principal components in  $m-\tau$  space) and  $U_n$  are uncorrelated factor loadings.

- Each eigen-surface corresponds to a ‘mode of deformation’ of the IV surface: level, slope (skew), curvature (smile thickness),
- Mean-reverting factor model
  - After finding the eigenmodes ( $f_k$ ), they model the factor loadings ( $x_k(t)$ ) (projections of the IV surface onto these modes) as mean-reverting Ornstein–Uhlenbeck processes, driven by independent noise sources, which can be Wiener or jump processes
  - This gives a low-dimensional stochastic model of the entire IV surface
- Key practical features
  - Data, smoothing, and surface construction
    - They use end-of-day SP500 and FTSE index options over ~1-2 years, focusing on liquid out-of-the-money options with moneyness in [0.5, 1.5] and maturities from around 1 month to 1 year
    - For each day, they construct a smooth IV surface on a fixed grid using a Nadaraya–Watson kernel estimator (Gaussian kernel in  $m$  and  $\tau$  with data-driven bandwidths). This is crucial: the KL decomposition is done on smoothed surfaces, not raw noisy quotes.
  - Dynamics, correlations, and risk implications
    - They explicitly connect this to Vega risk: since multiple factors move (not just level, and not perfectly linked to the underlying), delta hedging is not enough; a factor-based Vega hedge is needed. They also show how their model extends “sticky delta” by adding stochastic deformations around the smile.
- Relevance
  - The paper provides canonical empirical stylized facts about IV surface dynamics:
    - low-dimensional factor structure (level / skew / curvature),
    - mean-reversion times around 1–2 months,
    - strong negative correlation of IV level with underlying (leverage),
    - relatively weak link between underlying and shape factors.
  - Justification for modelling implied vol surfaces directly
    - Our approach is training a diffusion model directly on forward curves + IV surfaces, so fits exactly this philosophy. We cite them as the classical

- justification for a “market-based” modelling approach rather than modelling the underlying S or instantaneous volatility only.
- Baseline generative model to compare against
  - Their factor-OU model is essentially a linear generative model for IV surfaces: simulate OU factors -> reconstruct surface via eigenmodes -> price options via Black–Scholes.

[2] Deep Learning from Implied Volatility Surfaces [Kelly, Bryan T. and Kuznetsov, Boris and Malamud, Semyon and Xu, Teng Andrea, 2023]

- Link
  - <https://dx.doi.org/10.2139/ssrn.4531181>
- Context
  - IV surface = image (moneyness  $\times$  maturity) containing rich info about state-contingent risk premia and return distribution.
  - Economic theory:
    - local derivatives of IV (Breeden–Litzenberger, Dupire) link surface geometry to Arrow–Debreu prices and volatility.
  - Empirical problem:
    - IV grid is discrete, noisy, illiquid -> hard to compute theory-driven local features directly.
- Theory part – key concepts
  - IV surface as structured image
    - 2D grid of IVs on (delta / moneyness, maturity); “universal local features” = non-linear functions of neighbouring pixels.
  - Economic link
    - Cross-moneyness derivatives  $\rightarrow$  Arrow–Debreu state prices (Breeden–Litzenberger).
    - Term-structure slope  $\rightarrow$  local variance via Dupire.
  - CNN inductive bias
    - locality, translation / rotation invariance
    - shared filters across the surface
    - better suited than fully-connected DNNs for structured IV data
  - Ensemble complexity
    - many local minima in CNN loss landscape
    - averaging many randomly initialized CNNs boosts performance (virtue of complexity)
  - Gradient outer product & principal linear features
    - new ML object; eigenvectors define “principal linear features” (PC-analogue)  $\rightarrow$  no linear feature sparsity; >100 linear features needed to explain predictive content of IV.

$$w_{*,t} = \arg \min_w \ell(w), \quad \ell(w) = \sum_{\theta=t-T}^t \sum_{i=1}^{N_\theta} (R_{i,\theta+1} - f(IV_{i,\theta}; w))^2,$$

- Practical part – what they actually do

- Data / preprocessing
  - Critical because we'll have to gather data for each day options based on some logic to keep sizing constant in our paper as well
  - OptionMetrics IvyDB IV surfaces → normalized grid (~10 maturities × 34 deltas) per stock-day; remove very short (10-day) expiry; handle missing / incomplete images.
- Models
  - CNN1, CNN4, CNN5 = convolutional nets with 1 / 4 / 5 conv layers; depth = complexity; complexity metric  $\sim \# \text{parameters} / \text{sample size}$ .

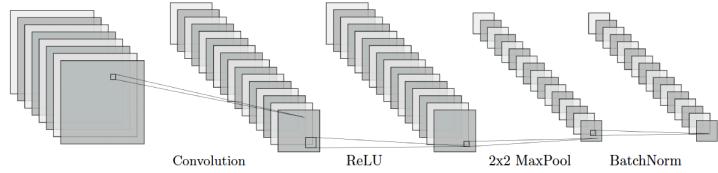


Figure 3: The figure above shows a building block of the CNN model consisting of a convolutional layer with a  $3 \times 3$  filter, a ReLU layer,  $2 \times 2$  max-pooling, and batch normalization layers. Note the max-pooling layer shrinks the height and width of the input by half and keeps the same depth.

- Ensembles of up to 100 randomly initialized CNNs
- also baselines: linear “kitchen-sink” ridge model and fully-connected NN1
- Main findings:
  - Deep CNN ensembles significantly predict 1-month stock returns using only month-end IV surface; out-of-sample Sharpe for H-L strategy rises from ~0.9 (single model) to ~2.7 (ensemble of 100 for deepest CNNs).
  - Alphas remain significant relative to many equity and option-based factors; robust to transaction costs and short-sale constraints.
  - Principal linear features: performance keeps improving as number of features  $P$  grows; need  $>100$  to capture predictive content → very high feature complexity

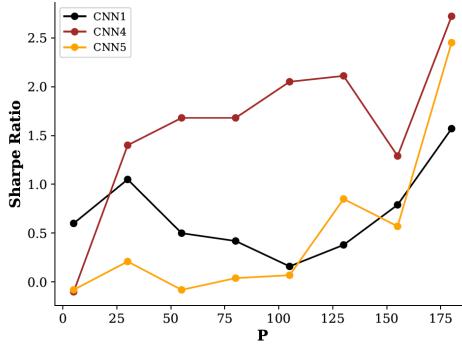


Figure 5: The figures above show the Sharpe Ratio of our H-L strategy (5) as a function of  $P$ , the number of principal features, based on the function  $f_P(x)$  constructed using Algorithm 1. The experiment is run separately for each of the CNN1, CNN4, and CNN5 models.

- Why this paper matters for our diffusion-IV project
  - Justification for modelling the full surface
    - shows that local geometry of the entire IV surface carries rich, non-linear predictive info that cannot be captured by a few summary statistics (level, slope, skew, convexity)
    - our conditional diffusion model over full surfaces is aligned with this “feature-rich” view
  - Structured-data perspective
    - treats IV as an image with locality and spatial structure; supports the idea of using structured generative models (diffusion with moneyness / maturity structure) rather than flat feature vectors.
  - Complexity vs parsimony
    - their evidence of ensemble and feature complexity provides a motivation to use flexible high-dimensional generative models
    - our work is going to base economic structure (no-arbitrage, P–Q split, forward curve consistency) on top of ML flexibility
  - Positioning
    - their CNNs are discriminative (return prediction);
    - our diffusion is generative (joint future paths of forward curve + IV surface with risk-management P&L metrics) – so we can explicitly present the model as complementary: instead of extracting predictive features, we will try to simulate economically consistent scenarios that could be fed into similar CNN-style predictors or used directly for hedging / risk.

[3] VolGAN: A Generative Model for Arbitrage-Free Implied Volatility Surfaces [Milena Vuletić and Rama Cont, 2024]

- Link
  - <https://www.tandfonline.com/doi/full/10.1080/1350486X.2025.2471317>
- Context
  - Conditional GAN for joint dynamics of underlying return + IV surface.
  - Trained on SPX options; aims at realistic scenarios + static no-arbitrage.
  - Used for forecasting, VIX simulation, and hedging option portfolios.
- Theory part
  - Static arbitrage penalty

Following Cont and Vuletic (2023), we define the *arbitrage penalty* associated with the (discretely sampled) volatility surface  $\sigma(\mathbf{m}, \boldsymbol{\tau})$  as:

$$\Phi(\sigma(\mathbf{m}, \boldsymbol{\tau})) = p_1(\sigma(\mathbf{m}, \boldsymbol{\tau})) + p_2(\sigma(\mathbf{m}, \boldsymbol{\tau})) + p_3(\sigma(\mathbf{m}, \boldsymbol{\tau})). \quad (2)$$

where the functions  $p_1, p_2, p_3$  measure violations of calendar, call and butterfly arbitrage constraints, respectively:

$$p_1(\sigma(\mathbf{m}, \boldsymbol{\tau})) = \sum_{i=1}^{N_m} \sum_{j=1}^{N_\tau} \left( \tau_j \frac{c(m_i, \tau_j) - c(m_i, \tau_{j+1})}{\tau_{j+1} - \tau_j} \right)^+, \quad (3)$$

$$p_2(\sigma(\mathbf{m}, \boldsymbol{\tau})) = \sum_{i=1}^{N_m} \sum_{j=1}^{N_\tau} \left( \frac{c(m_{i+1}, \tau_j) - c(m_i, \tau_j)}{m_{i+1} - m_i} \right)^+, \quad (4)$$

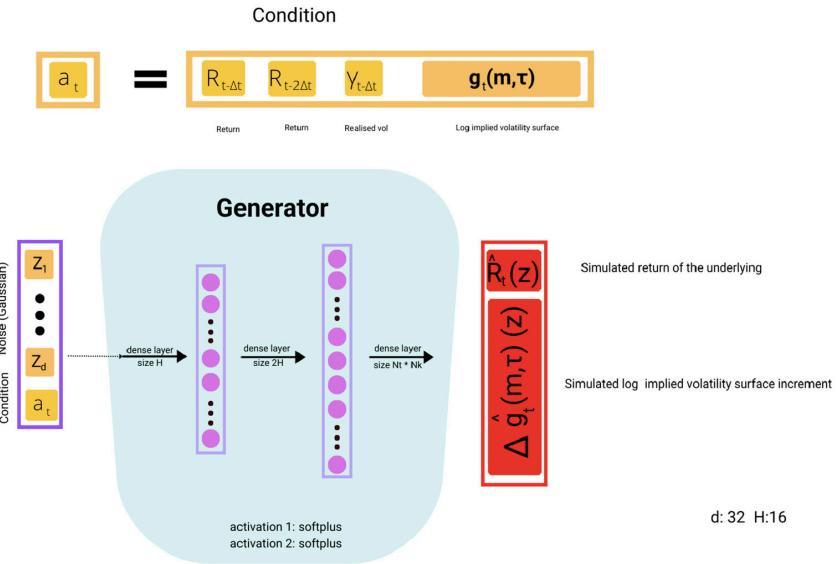
$$p_3(\sigma(\mathbf{m}, \boldsymbol{\tau})) = \sum_{i=1}^{N_m} \sum_{j=1}^{N_\tau} \left( \frac{c(m_i, \tau_j) - c(m_{i-1}, \tau_j)}{m_i - m_{i-1}} - \frac{c(m_{i+1}, \tau_j) - c(m_i, \tau_j)}{m_{i+1} - m_i} \right)^+. \quad (5)$$

Static arbitrage constraints (Davis and Hobson 2007) are then equivalent to

$$\Phi(\sigma(\mathbf{m}, \boldsymbol{\tau})) = 0$$

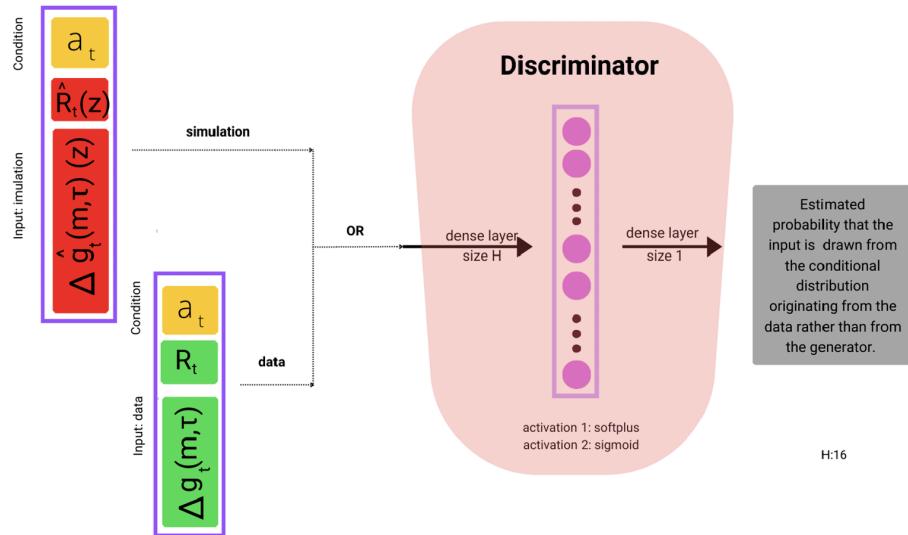
and the magnitude of  $\Phi(\sigma(\mathbf{m}, \boldsymbol{\tau}))$  can be considered as a ‘distance’ from the set of arbitrage-free implied volatility surfaces.

- Shape constraints for IV surface
  - Monotonicity in maturity ( $\partial \tau C \geq 0$ ), decreasing in moneyness ( $\partial m C \leq 0$ ), convex in moneyness ( $\partial^2 m C \geq 0$ ).
  - Translated into constraints on  $\sigma(\mathbf{m}, \boldsymbol{\tau})$  and its derivatives
- Conditional GAN architecture (VolGAN)
  - Condition vector
    - previous IV surface  $g^\square(\mathbf{m}, \boldsymbol{\tau})$
    - last two log-returns
    - previous realized vol
  - Generator  $G(a^\square, z) \rightarrow (\text{next log-return}, \log\text{-IV increment } \Delta g^\square(\mathbf{m}, \boldsymbol{\tau}))$



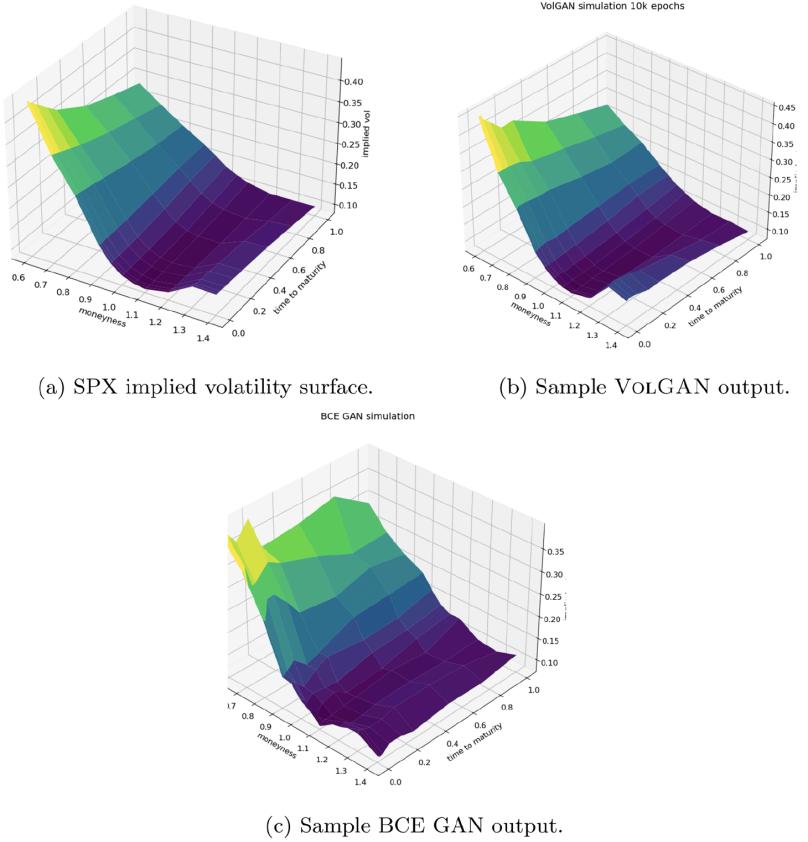
**Figure 1.** VoLGAN generator architecture.

- Discriminator  $D(a_\square, (R, \Delta g))$  distinguishes real vs generated



**Figure 2.** VoLGAN discriminator architecture.

- Smoothness penalty (Sobolev seminorm)
  - $L_\square$  and  $L_\tau$  penalties on discrete derivatives in  $m$  and  $\tau$  for log-IV surface; encourage smooth, “PDE-like” surfaces
  - The result of their work is presented below - we can see it is super smooth, so we have to keep that discovery in our work as well



**Figure 6.** Implied volatility surfaces generated using (b) VolGAN (c) classical GAN, compared with (a) SPX implied volatility surface.

- Empirical part
  - Hedging application
    - Scenario-based regression hedging of a 1-month straddle using VolGAN scenarios
    - Compare BS delta, BS delta-vega, and VolGAN + regression (LASSO / ATM only)
  - Relevance to our conditional diffusion paper
    - Same object, different generator
      - VolGAN provides a GAN-based baseline for our diffusion-based generator.
    - Arbitrage handling paradigm:
      - VolGAN uses smoothness penalties + scenario re-weighting based on a discrete arbitrage functional  $\phi(\sigma)$ .
      - We can borrow their arbitrage penalty functional and shape constraints
    - Conditional structure & inputs:
      - Their conditioning set (past IV surface + returns + realized vol) is a good template for our conditioning variables in a diffusion model (we just extend it with forward/futures curve, extra history, etc.).
    - Our contribution
      - replace GAN with conditional diffusion
      - add joint forward curve + IV surface, more explicit P/Q separation
      - emphasise P&L-based evaluation using a similar hedging setup

[4] Controllable Generation of Implied Volatility Surfaces with Variational Autoencoders  
 [Wang, Liu, Vuik, 2025]

- Link
  - <https://arxiv.org/abs/2509.01743>
- Context
  - Proposes a controllable VAE that generates IV surfaces with specified shape features (level, slope, curvature, term-structure slope).
  - Adds post-generation arbitrage repair in latent space (calendar + butterfly).
- Experiment features
  - Dataset
    - 60k synthetic IVSs on fixed  $28 \times 28$  grid in  $(m, \tau)$ , from Heston + SABR parameter ranges
    - VAE: latent dim  $z=5$ , ResNet encoder/decoder,  $\beta$  tuned per experiment

Table 1: Hyperparameters used in the controllable VAE model.

Hyperparameter	Setting
Input Dimension ( $x$ )	784
Control Dimension ( $y$ )	problem-specific
Latent Dimension ( $z$ )	5
Encoder hidden layers	[256, 128]
Decoder hidden layers	[128, 256]
Batch size	64
Activation function	ReLU
Optimizer	Adam
Learning rate	$3 \times 10^{-4}$
$\beta$	problem-specific
Max epochs	5000

- Relevance
  - Their VAE is a static generator with feature knobs; the diffusion we are gonna use is a dynamic conditional generator (next-day surface + forward curve).
  - Here we can position our model as: same economic features + richer temporal dynamics and joint forward/IV structure.
  - Arbitrage treatment template:
    - They use post-hoc latent optimization with explicit calendar/butterfly penalties.
    - Here we can adopt their penalty functions and diagnostics as evaluation tools;
    - In the contrary, our diffusion aims to build no-arbitrage into the generative dynamics, not only fix outputs afterwards.

[5] Diffusion-Based Generative Modeling of Financial Time Series [2025]

- <https://hdl.handle.net/10012/22497>
- Problem: generative modeling of multi-asset financial time series beyond GBM/Heston-type parametrics.
- Proposed method: Elucidated Diffusion Model + NCSN++ backbone adapted to returns, with “Ambient Diffusion” variance correction and analytic noise schedule.

- Data/metrics: synthetic GBM/Heston/Merton + real SPY, NVDA, BTC etc; evaluated on distribution fit, SDE parameter recovery, option pricing and risk metrics (VaR/CVaR).
- Result: ambient-corrected diffusion markedly reduces volatility bias and pricing error versus vanilla EDM across assets.

[6] Volatility Surface Completion using Score-Based Generative Models [2023]

- Problem: “volatility completion” – filling missing IV quotes on a strike–maturity grid while enforcing static no-arbitrage.
- Proposed method: treat IV surface as an image on a fixed grid and use a noise-conditional score network with Langevin inpainting, modified to impose butterfly + calendar constraints during sampling.
- Data and metrics: Heston-generated IVS on  $8 \times 8$  and  $16 \times 16$  grids; interpolation and randomized masks up to 80% missing.
- Result: interpolation errors  $\approx 10^{-4}$  and max error <0.5% even with 80% missing, while keeping surfaces essentially arbitrage-free.

[7] Forecasting Implied Volatility Surface with Generative Diffusion Models [2025]

- <https://arxiv.org/abs/2511.07571>
- Problem: 1-day-ahead forecasting of arbitrage-free IV surfaces on SPX using a data-driven model that handles path dependence and residual arbitrage in training data.
- Proposed method: conditional DDPM on a  $9 \times 9$  (moneyness, maturity) grid, conditioned on EWMA of past IV surfaces, returns/squared returns, and VIX; loss combines reconstruction with SNR-weighted arbitrage penalty.
- Key theory: shows the arbitrage penalty introduces a small, controllable bias while steering the model toward the arbitrage-free manifold.
- Experimental result: on SPX OTM options, diffusion model beats VolGAN variants on MAPE and delivers better-calibrated 90% CIs for vols across ATM/OTM/ITM buckets and horizons.

[8] Generating the Term Structure of Interest Rates with Diffusion Models [2025]

- [https://papers.ssrn.com/sol3/papers.cfm?abstract\\_id=5493026](https://papers.ssrn.com/sol3/papers.cfm?abstract_id=5493026)
- Problem: generative modeling of zero-coupon yield curves (term structures) conditional on macro/market variables across currencies.
- Proposed method: conditional DDPM with v-parameterization and cross-attention U-Net, generating dense OIS-based yield curves for JPY/USD/GBP given macro and rate inputs.
- Variants: direct curve generation vs generating day-to-day differences; plus a faster Nelson-Siegel-Svensson (NSS) factor-based diffusion.
- Results: realistic yield-curve shapes and dynamics, good 6-month out-of-sample behavior; NSS factor model cuts training/inference time by  $\approx 40\text{--}44\%$  with similar quality.

[9] A Neural Network Approach to Understanding Implied Volatility Movements [2019]

- [https://www.researchgate.net/publication/341383844\\_A\\_neural\\_network\\_approach\\_to\\_understanding\\_implied\\_volatility\\_movements](https://www.researchgate.net/publication/341383844_A_neural_network_approach_to_understanding_implied_volatility_movements)
- Problem: empirical modeling of how the IV surface moves with index returns, moneyness and maturity for S&P 500 options.
- Proposed method: “three-feature” NN (return, delta-moneyness, time to maturity) and “four-feature” NN (adds VIX) trained on ~2M daily SPX call observations (2010–2017).
- Use cases: compare empirical surface dynamics to stochastic-vol models; compute minimum-variance delta that accounts for expected IV changes.
- Key result: NNs significantly outperform a simple analytic regression model across regimes; VIX as fourth feature further improves fit, especially in high-vol markets.